

final-P1-sol

December 19, 2020

1 Academic Integrity Statement

As a matter of Departmental policy, **we are required to give you a 0** unless you **type your name** after the following statement:

I certify on my honor that I have neither given nor received any help, or used any non-permitted resources, while completing this evaluation.

[TYPE YOUR NAME HERE]

2 Problem 1 (50 points)

Rampant disinformation—often called “fake news”—has emerged as one of the fundamental crises over our time.

There is a growing movement for online platforms to regulate fake news. Doing so at scale requires combing through millions of news items every day, making it very expensive to do by hand. Can an algorithm do it instead?

The following two URLs each contain part of a data set.

- **Fake news items:** <https://raw.githubusercontent.com/PhilChodrow/PIC16A/master/datasets/fake>
- **Real news items:** <https://raw.githubusercontent.com/PhilChodrow/PIC16A/master/datasets/fake>

Use the data at these urls to **construct a fake news classifier**.

1. Your model must be able to **make predictions** about whether or not an unseen news item is fake or real.
2. Because fake news models must be able to make millions of predictions per day, it must be able to make predictions very quickly. More columns mean more computation time. **Your final model should use no more than 50 columns.**

You are free to create any columns that you need, and to use any functions that we have or have not covered in the course. You may also use any machine learning model.

Please use Markdown headers with `##` signs to clearly distinguish the different stages of your solution.

2.0.1 Requirements

1. Any operations that you perform multiple times (such processing that you perform on both the training and test sets) must be contained in function with informative docstrings. Comments and explanations are expected throughout. It is especially important to explain how you chose the columns to use in your final model.
2. You should not use for-loops to iterate over the rows of data frames or arrays.
3. You must fit your model on the training data, and not use the test data for fitting.

2.0.2 Hints

- `pd.concat()` is a good way to combine data frames.
- Try fitting a model with as many columns as you want first. See if you can get a representation of which columns are important, and then select your final columns from this list.
- In class, we talked about greedy stagewise feature selection and exhaustive enumeration for determining a good set of columns. Neither of these methods are recommended for this problem.
- If you want to be creative about your model choice, then please go for it. If you want a safe option, try logistic regression.
- If a model takes too long to fit on the full data set, try fitting it on, say, 10% of the data.
- You might find the some of the [cheatsheets](#) to be helpful.

2.0.3 Rubric

- **(15 points)**: clearly written code that makes economical use of skills from the course to manipulate data.
- **(15 points)**: comments, explanatory surrounding text, and docstrings for any functions and classes.
- **(20 points)**: computed according to the formula $20 \times \text{score}$, where `score` is your model's prediction performance on unseen data. Models that use more than 50 columns can receive up to 15 of these points. Scores will be rounded up. For example, if you obtain an 84% predictive performance with 50 columns, then the score of $20 \times 0.84 = 16.8$ will be rounded up to 17 points.

3 My Solution

***Note:** many of you chose very different approaches here. Provided that you explained yourself, these generally received full credit. Some of you noticed that I made a silly mistake: the `subject` column is actually perfectly correlated with whether or not a news item was real, and so by using it you could get 100%. I didn't notice this and should have removed it! However, noticing something like this takes some work and some instincts, so solutions using this column still received full credit.*

First I'll download the data.

Solutions using the older multistep code shown in some lectures were also fine, but if you used it multiple times (rather than wrapping it in a function which you called twice), I made a small

deduction under the first rubric item.

```
[1]: import pandas as pd
fake_df = pd.read_csv("https://raw.githubusercontent.com/PhilChodrow/PIC16A/
↳master/datasets/fake_news/Fake.csv")
true_df = pd.read_csv("https://raw.githubusercontent.com/PhilChodrow/PIC16A/
↳master/datasets/fake_news/true.csv")
```

Next, I'll make indicators in each of the two pieces of the data indicating whether or not they are fake, and then combine them.

Some of you had problems from duplicated indices. Either passing `ignore_index` below or using `reset_index` later can solve these issues.

```
[2]: fake_df['fake_news'] = 1
true_df['fake_news'] = 0
df = pd.concat((fake_df, true_df), axis = 0, ignore_index=True)
df = df.drop(["title", "date", "subject"], axis = 1)
```

3.0.1 Prepare Data

Now I'll prepare my data. First, I'll extract the term-document matrix using only relatively common words, and add it to my data frame. Then, I'll create predictor and target variables.

```
[3]: from sklearn.feature_extraction.text import CountVectorizer

def prepare_data(df):
    """
    Prepare the fake news data set, passed as df, for subsequent machine_
    ↳learning tasks.

    1. Compute a term document matrix using sane defaults and add it to the df
    2. Prepare predictor X and target y.

    X is obtained by dropping the fake_news and complete text columns from df
    It therefore contains only the term-document matrix
    y is the fake_news column

    return: X, y, the prepared predictor and target data.
    """
    # 1. term-document matrix
    vec = CountVectorizer(max_df=.5, min_df = 1000, stop_words='english')
    counts = vec.fit_transform(df['text'])
    counts = counts.toarray()
    count_df = pd.DataFrame(counts, columns=vec.get_feature_names())

    # add to data frame
    out = pd.concat((df, count_df), axis = 1)
```

```
# 2. prepare predictor and target variables
X      = out.drop(['fake_news', 'text'], axis = 1)
y      = out['fake_news']

return X, y
```

```
[4]: X, y = prepare_data(df)
```

3.0.2 Train-Test Split

```
[5]: # for reproducibility
import numpy as np
np.random.seed(1234)

from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size = 0.5)
```

3.0.3 Identify Good Columns

So, I'm going to use logistic regression on the entire data set first. Then, I'll examine the coefficients, and use these to select the final 50 columns that I'm going to use.

First we'll create and train the model:

```
[6]: from sklearn.linear_model import LogisticRegression
LR = LogisticRegression(solver = "liblinear")
LR.fit(X_train, y_train)
```

```
[6]: LogisticRegression(solver='liblinear')
```

Now we'll evaluate. Because the model has so many columns to work with, the training score especially is very high! But the test score also isn't bad.

One could try CV here to improve even more, but I think this is good enough for now.

```
[7]: LR.score(X_train, y_train), LR.score(X_test, y_test)
```

```
[7]: (0.9952336406966903, 0.9656554857677402)
```

My approach to figuring out the most relevant words is to grab the top 50 words according to the absolute value of their coefficient.

```
[8]: # create data frame with the coefficients and words
result_df = pd.DataFrame({"coef" : LR.coef_[0], "word" : X_train.columns})
# compute the absolute value of each coefficient
result_df['abs'] = np.abs(result_df['coef'])
```

```
# sort by the absolute value and grab the best 50
top_words = result_df.sort_values('abs', ascending = False).head(50)
# display
top_words
```

```
[8]:
```

	coef	word	abs
502	4.283595	featured	4.283595
570	3.784056	gop	3.784056
559	3.663105	getty	3.663105
1051	3.617180	read	3.617180
1089	3.103784	rep	3.103784
945	2.957214	pic	2.957214
628	2.680299	https	2.680299
1119	-2.644479	reuters	2.644479
194	2.642317	breitbart	2.642317
1316	-2.593565	thursday	2.593565
1402	-2.465560	wednesday	2.465560
359	2.300493	daily	2.300493
1439	2.160724	yesterday	2.160724
1305	2.108230	thanks	2.108230
27	2.103557	21wire	2.103557
842	2.062048	morning	2.062048
637	2.051147	image	2.051147
26	2.039335	21st	2.039335
455	2.027216	entire	2.027216
260	-2.019130	citing	2.019130
541	-2.015438	friday	2.015438
838	-1.922789	monday	1.922789
1298	1.912387	terror	1.912387
1067	-1.878186	referring	1.878186
110	1.867583	apparently	1.867583
683	1.866951	isn	1.866951
877	-1.864864	nov	1.864864
122	1.853634	aren	1.853634
1432	1.824422	wouldn	1.824422
265	1.822876	claimed	1.822876
1395	1.807508	watch	1.807508
1100	-1.799013	representatives	1.799013
1181	-1.743681	sept	1.743681
278	1.734637	com	1.734637
1358	-1.704565	urged	1.704565
1172	1.703769	sen	1.703769
848	1.700586	mr	1.700586
1120	1.689319	revealed	1.689319
322	1.674222	controversial	1.674222
1293	1.671209	telling	1.671209
784	1.639280	lying	1.639280

1342	-1.639083	tuesday	1.639083
181	-1.636248	bit	1.636248
1357	1.614529	unless	1.614529
770	-1.598257	london	1.598257
402	1.597520	didn	1.597520
1147	-1.560926	saturday	1.560926
114	1.556755	appears	1.556755
614	1.516468	hill	1.516468
1248	1.515818	stated	1.515818

Now I'll just get the words on their own:

Typing/copy-pasting the words from your model isn't great practice (not reproducible) and led to a deduction under code quality.

```
[9]: top_words = top_words['word']
```

3.0.4 Final Model

Now I'll fit another logistic classifier using just the top words, and score it:

possible to do even better with other models!

```
[10]: LR2 = LogisticRegression()
      LR2.fit(X_train[top_words], y_train)
      LR2.score(X_test[top_words], y_test)

      # score on this problem would be 19/20
```

```
[10]: 0.9306873357387857
```